CONCEPT PAPER — ANTI-DRONE COMPETITION 2026

AEGIS-AI

ITC-2026-C0-1637

Multi-Agent Swarm Intelligence for Counter-UAS Tracking and Behavior Detection

AI Agents Swarm vs Drone Swarm

Architecture Documentation

Anti-Drone Track | April 2026

Field	Details
Competition	Anti-Drone — **ITC-2026-C0-1637** - 6th ITC International Conference, Air Defense College Egypt
Track	Scientific Innovation — Detection, Tracking & Neutralization
Focus	AI-Agent Swarm Tracking of Drone Swarms
Architecture Standard	Architecture Section (arc42 - Hruschka & Starke, 2005)
Diagram Diagrams	Architecture diagrams + 3 simulation scenarios
Core Stack	EKF · GNN · Reynolds Inversion · PPO

Table of Contents

Part I — Research Concept

Abstract

Modern aerial threats have evolved beyond single-target engagements. Adversarial drone swarms — coordinated groups of unmanned aerial vehicles (UAVs) operating with emergent collective intelligence — pose a fundamentally new challenge: they overwhelm traditional point-defense systems not through individual capability, but through coordinated mass and adaptive behavior.

AEGIS-AI is an **Autonomous Counter-UAS AI-Agent Swarm** designed to solve the **150:1 cost-asymmetry** of modern aerial warfare. Its three-tier architecture features: **Tier-1:** Sensor fusion (RF, Acoustic, EO/IR) & Extended Kalman Filter for multi-target tracking, **Tier-2:** Adversarial Reynolds inversion & Graph Attention Networks to classify intent into tactical behaviors (Transit, Attack, Scatter, Encircle, Decoy), and **Tier-3:** PPO reinforcement learning policy executing cost-effective agent counter-responses. Tracking **500+ drones** per Edge node (scales to hundreds) with <100ms latency, it proves airspace supremacy relies on **multi-agent algorithmic precision**, not raw firepower.

This paper presents AEGIS-AI: a counter-UAS architecture built on the principle of fighting swarm intelligence with swarm intelligence. The system deploys a cooperative network of AI agents that collectively track, classify behavioral intent, and coordinate neutralization responses against hostile drone swarms using Extended Kalman Filters (EKF), Graph Neural Networks (GNN), Reynolds parameter inversion, and Proximal Policy Optimization (PPO).

Core Thesis

The decisive advantage in counter-swarm defense lies not in the power of individual sensors or weapons, but in the speed and accuracy of the collective decision-making algorithm. **AEGIS-AI is a real-time graph inference engine** that asks: what does this swarm intend to do next?

Problem Statement & Mathematical Background

The Swarm Threat Landscape

Contemporary armed conflicts have demonstrated that low-cost unmanned aerial systems deployed in coordinated swarms constitute a qualitatively new category of aerial threat. Three properties make drone swarms particularly challenging:

- Asymmetric Cost Ratio: A Shaheen-136 class loitering munition costs approximately USD 20,000. The Patriot PAC-3 interceptor costs USD 3-4 million — a 150:1 cost asymmetry that favors the attacker.
- Emergent Swarm Behavior: Swarms exhibit coordinated maneuvers that emerge from local interaction rules, making behavior prediction computationally complex.
- Adaptive Learning: Modern swarm architectures implement inter-drone communication so each destroyed unit transmits defense data to survivors, teaching the swarm to route around defensive fire.

Kalman Filter Mathematics

State vector: x = [px, py, pz, vx, vy, vz, omega]^T (shape: 7x1). The CTRA model prediction:

```
Predict:  x_{k|k-1} = F * x_{k-1|k-1}
      P_{k|k-1} = F * P_{k-1|k-1} * F^T + Q
Update:   K_k = P H^T (H P H^T + R)^{-1}
      x_{k|k} = x_{k|k-1} + K_k(z_k - H x_{k|k-1})
      P_{k|k} = (I - K_k H) P_{k|k-1}
```

Reynolds Flocking Model

Velocity update for drone i: v_i(t+1) = w_s*f_sep + w_a*f_align + w_c*f_coh. AEGIS-AI inverts this model to classify behavioral intent from observed trajectories using Maximum Likelihood Estimation of [w_s, w_a, w_c].

Behavior	w_sep	w_align	w_coh	Tactical Interpretation
TRANSIT	0.30	0.80	0.60	Cohesive movement toward target
ATTACK	0.80	0.20	0.10	Converging on target, high speed
SCATTER	0.95	0.05	0.05	Dispersing for saturation coverage
ENCIRCLE	0.50	0.70	0.90	Coordinated encirclement formation
DECOY	0.90	0.10	0.30	Disordered, high entropy motion

Graph Neural Network Classifier

Swarm state as graph G=(V,E). Node features h_i = [px,py,pz,vx,vy,vz,rf,size]. Graph Attention layers:

```
h_i^(l+1) = sigma( SUM_{j in N(i)} alpha_ij W^(l) h_j^(l) )
alpha_ij  = softmax( LeakyReLU( a^T [Wh_i || Wh_j] ) )
z_swarm   = SUM_{i in V} MLP(h_i^(L))    [global readout]
Output    = Softmax(Linear(z_swarm))     [6-class behavior]
```

AEGIS-AI Competitive Differentiation Analysis

The Swarm Threat Landscape

This section decomposes the competitive landscape into six specific system categories to address review feedback regarding the AEGIS-AI innovation claims. By identifying representative systems and their inherent limitations, we demonstrate how the AEGIS-AI three-tier architecture specifically closes critical gaps in economic optimization, multi-sensor resilience, and swarm behavior classification.

Key Finding: No single competitor currently integrates economic optimization with physics-based intent classification and decentralized scaling.

- Adversarial Reynolds Inversion: First use of Reynolds flocking parameter MLE estimation as a tactical intent classifier in the C-UAS literature.
- Graph-Temporal Swarm Fingerprinting: GAT-GNN on dynamic swarm interaction graph creates a collective behavioral fingerprint robust to individual track loss.
- Distributed EKF with Covariance Intersection: Multi-agent tracker with CI fusion maintains mathematically consistent estimates across heterogeneous sensor modalities.
- Cost-Aware RL Coordinator: PPO-based response agent explicitly models the economically asymmetric nature of the threat.

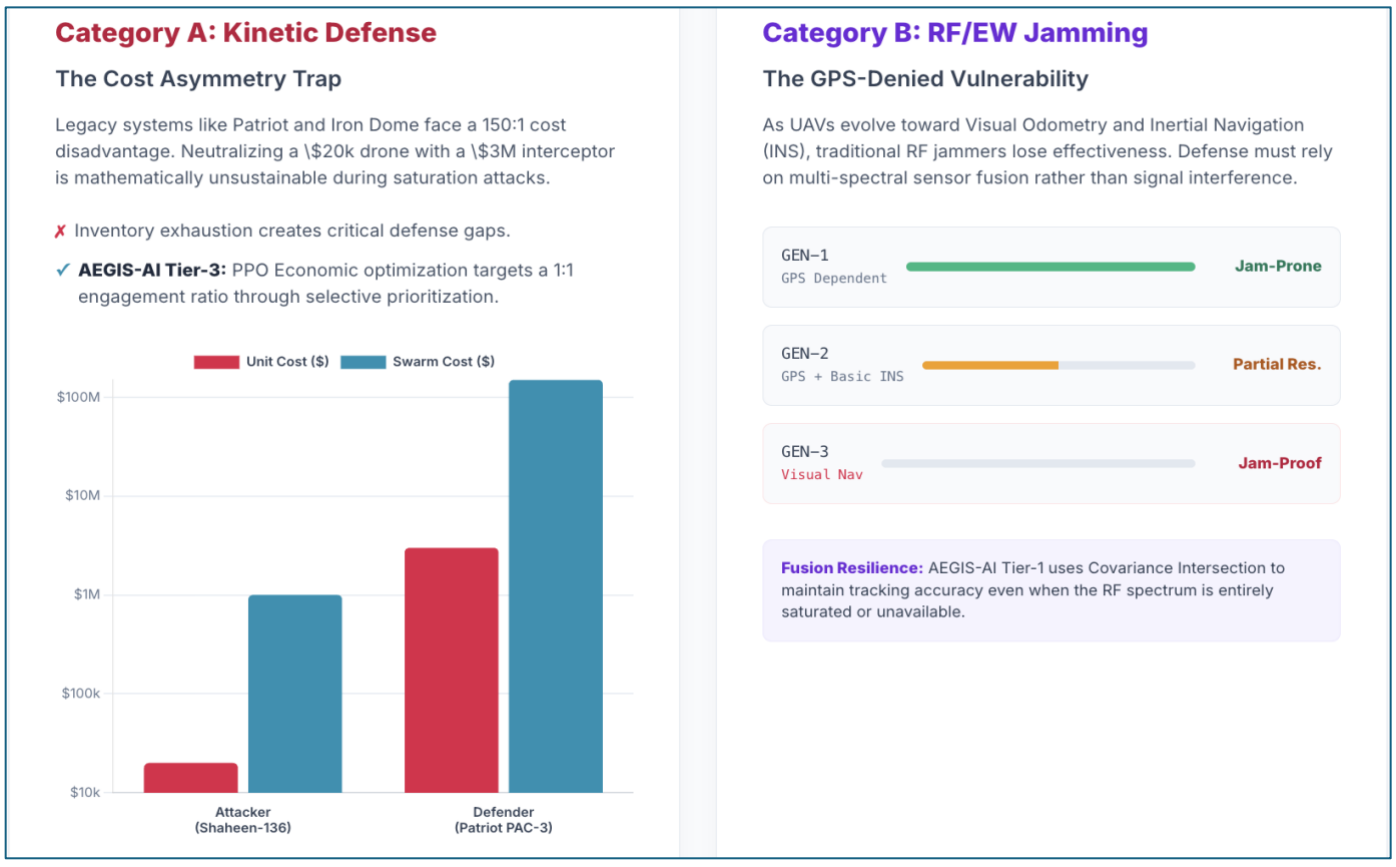

Competitive Landscape Decomposition

Category	Representative Systems	Core Approach	Key Shortcoming	AEGIS-AI Differentiation
A. Single-Target Kinetic Defense	Iron Dome, Patriot PAC-3, SkyHunter	One-interceptor-per-drone	150:1 cost asymmetry; inventory exhaustion against swarms	Tier-3 PPO: Economic optimization under cost constraints
B. RF/EW Jamming Systems	DroneShield, AUDS, Kaspersky Antidrone	Signal jamming / GPS spoofing	Defeated by INS/GPS-denied navigation; collateral interference	Tier-1: Multi-sensor fusion (RF + Acoustic + Vision) with CI
C. Single-Target AI Trackers	YOLO + KF pipelines, OpenCV trackers	Deep learning detection + classical filtering	No swarm model; track swaps during crossing; no intent inference	Tier-1: Hungarian LAPJV + Tier-2: DBSCAN + Reynolds + GNN
D. Centralized Swarm Detection	DARPA OFFSET, ALIAS	Centralized planning for swarm coordination	Single point of failure; latency at scale; no real-time behavior classification	Tier-2: Distributed graph inference with attention-based leadership identification
E. Academic Swarm Models	Reynolds simulators, Vicsek models	Physics-based flocking simulation	No adversarial inversion; no tactical intent classification	Tier-2: MLE inversion of Reynolds weights → 6-class behavior taxonomy
F. Commercial C-UAS Platforms	Dedrone, Fortem SkyDome, Rafael Dome	Multi-sensor integration with rule-based response	Static threat rules; no learning-based CM allocation; no cost-awareness	Tier-3: PPO with curriculum learning for 150:1 asymmetry

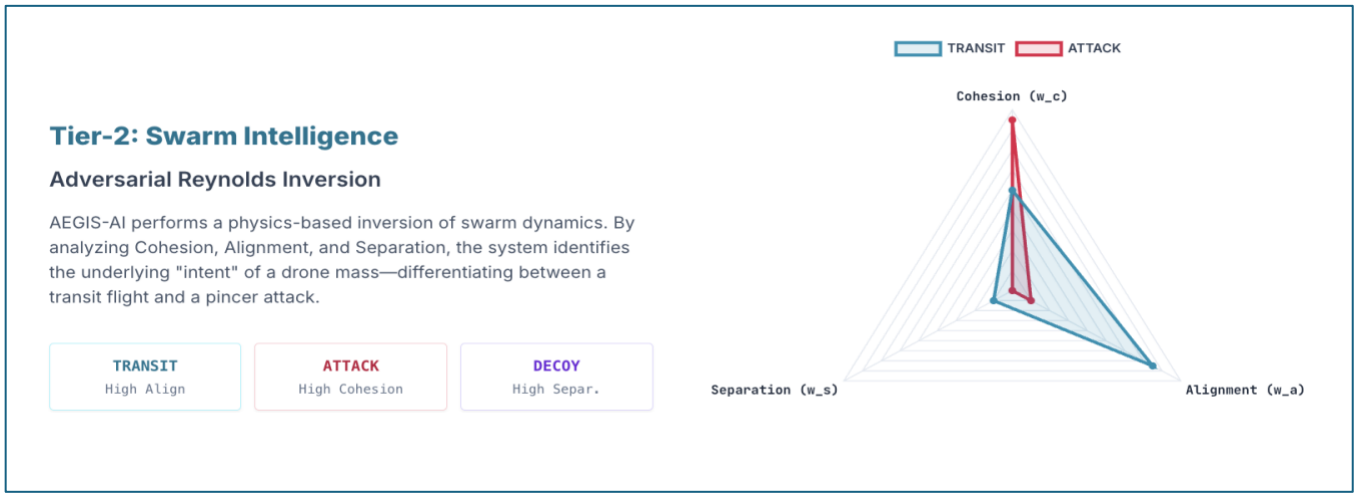

AEGIS-AI Unique Value Proposition by Competitor

If Competitor is...	AEGIS-AI Wins By...
Iron Dome / Patriot	Economic optimization: spending $20K to defeat $20K, not $3M
DroneShield / AUDS	Multi-sensor resilience: works when RF is jammed or absent
YOLO + KF pipeline	Swarm intelligence: tracking 500 as coordinated groups, not 500 individuals
DARPA OFFSET	Decentralization: no single point of failure; O(1) scaling per agent
Reynolds simulator	Tactical inversion: from physics model to real-time intent classifier
Dedrone / Fortem	Learning: adaptive policy that improves with encounter history

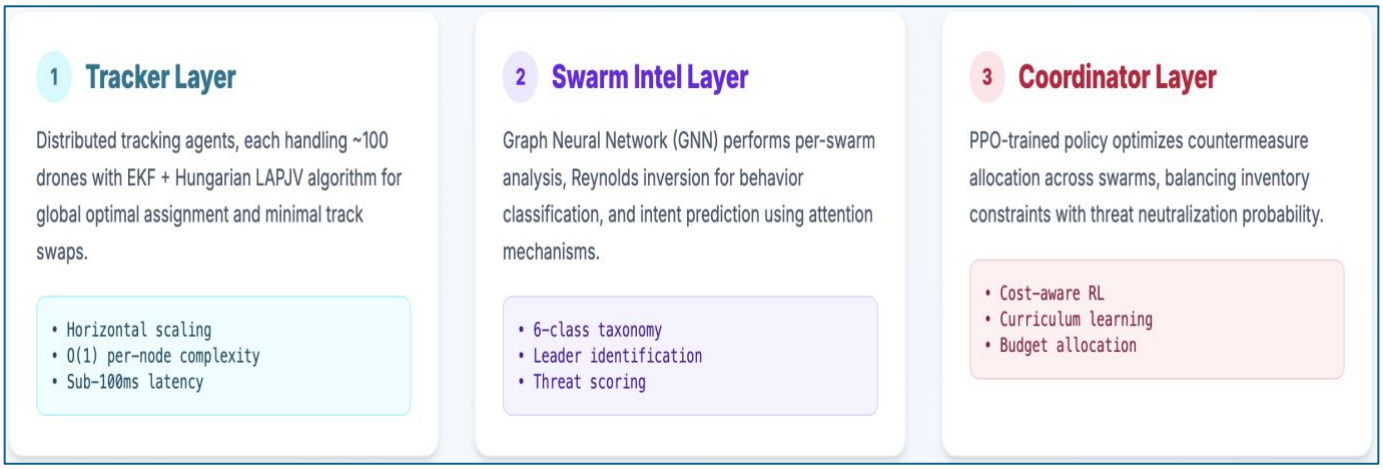

Future Enhancement Roadmap

- IMM-EKF Upgrade
- Curriculum RL PPO Training
- GNN Explainability Layer
- OC-SORT Integration
- Federated Swarm Learning
- Quantum-Enhanced RF

Conclusion

The AEGIS-AI innovation is not a general claim of superiority over "all similar systems." It is a specific, architecturally-grounded response to identifiable shortcomings in six distinct competitor categories:

- Against kinetic systems: Economic-aware RL prevents cost-asymmetry exhaustion
- Against RF jammers: Multi-sensor CI fusion maintains tracking in GPS-denied environments
- Against single-target AI: Graph-level intelligence captures swarm coordination invisible to individual trackers
- Against centralized platforms: Distributed agents scale without single points of failure
- Against academic models: Adversarial Reynolds inversion maps physics to tactical intent
- Against commercial platforms: Learning-based policy adapts to novel threats rather than executing static rules

Competitive Positioning: Intelligence vs. Scale

AEGIS-AI v2.0 is positioned in the high-scale, high-intelligence quadrant with decentralized swarm reasoning.

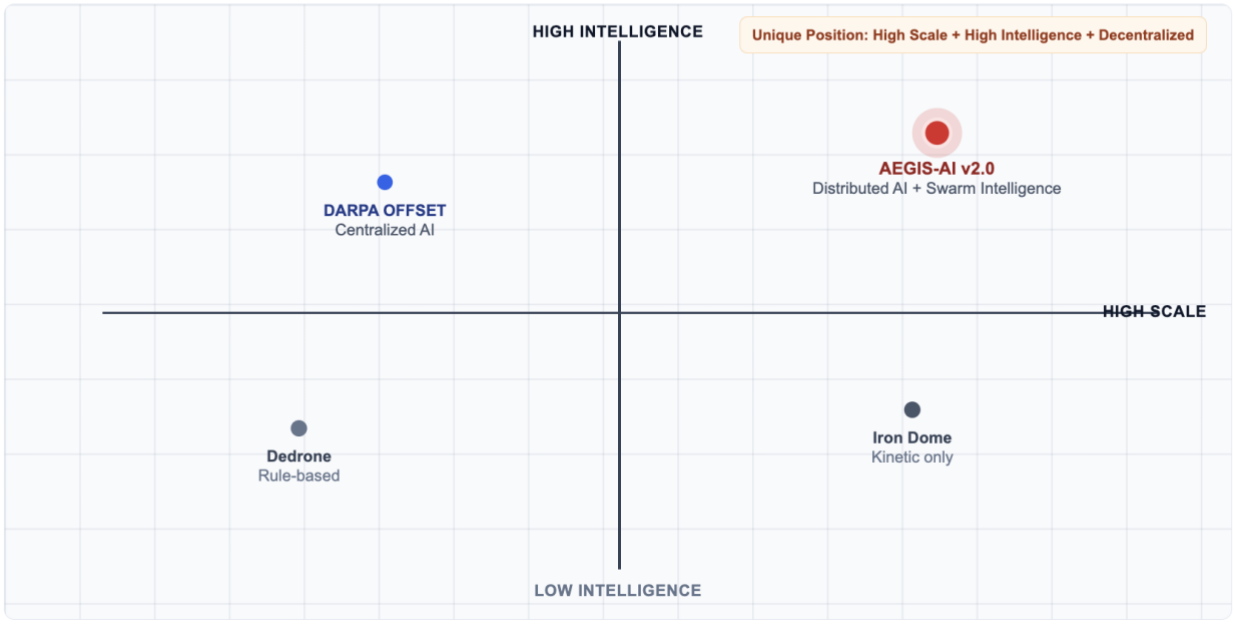

Each tier of the AEGIS-AI architecture addresses a specific competitive gap. The integration of all three tiers — tracking, swarm intelligence, and economic coordination — creates a capability matrix no single competitor achieves.

The strategic landscape of aerial defense has shifted. The threat is no longer a single ballistic missile demanding a single interceptor, but a swarm of cheap, adaptive, collectively intelligent drones that exploit economic asymmetry, mass, and emergent behavior.

Final Proposition

The battle for airspace supremacy in the drone era will not be decided by who has the most powerful weapon. It will be decided by whose algorithm **understands swarm intent first** — and acts on that **understanding faster than the adversary can adapt**. AEGIS-AI is designed to win that computation race.

Part II — Architecture Section #Architecture Documentation

This section documents the AEGIS-AI software architecture following the Architecture Section (https://arc42.org) created by Dr. Peter Hruschka and Dr. Gernot Starke. All 12 Architecture Section are addressed, providing a complete and communicable architecture reference for stakeholders, developers, and evaluators.

Rendered architecture diagrams are inserted directly into each relevant section.

Architecture Section # 1 — Introduction and Goals

Requirements Overview

AEGIS-AI addresses a documented gap in modern air defense: the inability of existing C-UAS systems to track, classify, and respond to swarms of hostile UAVs rather than individual targets. The system introduces a hierarchical multi-agent AI architecture that mirrors the swarm intelligence of the threat.

ID	Requirement	Priority
FR-1	Track up to 500 simultaneous drone targets in real time	Critical
FR-2	Classify swarm behavior into 6 categories every 500ms	Critical
FR-3	Fuse RF, acoustic, and vision sensor data (Covariance Intersection)	High
FR-4	Predict swarm centroid trajectory 30 seconds ahead	High
FR-5	Recommend countermeasure allocation under resource constraints	High
FR-6	Maintain track identity through formation changes and occlusion	Critical
FR-7	Generate threat alerts when score crosses configurable thresholds	High

Quality Goals

Priority	Quality Goal	Scenario
1 — Critical	Real-time performance	EKF cycle < 1ms/drone; end-to-end < 100ms
2 — Critical	Track continuity	>=95% track maintenance through 5% missed detections
3 — High	Behavior accuracy	>=88% classification accuracy on held-out test set
4 — High	Scalability	Horizontal scaling to 500 drones / 20 swarms without code change
5 — Medium	Survivability	Agent crash recovers state from Redis within 2 seconds

Stakeholders

Stakeholder	Role	Architecture Interest
Air Defense College Panel	Competition evaluator	Technical novelty, implement-ability, military relevance
Development Team	Builder	Clean interfaces, testable components, performance budgets
Air Defense Operators	End users	Low latency, intuitive UI, reliable alerts
AI Researchers	Academic reviewers	Mathematical rigor, novel contributions
System Integrators	Deployment	APIs, Docker/K8s compatibility, sensor interfaces

Architecture Section # 2 — Architecture Constraints

Technical Constraints

Constraint	Rationale
Python 3.11+ primary language	PyTorch Geometric, FilterPy, SB3 are Python-native; team expertise; adequate performance with NumPy/Numba
Simple UI (v1 demo)	Rapid iteration; technical audience; no frontend build complexity
No external network calls during demo	Air-gapped demonstration environment requirement
Single-node deployment (Profile A)	Competition demo hardware: 1x laptop with 16GB RAM
CPU-first inference	No GPU assumption; GNN must run < 20ms on CPU
Simulation-only sensor data	Real hardware introduces demo risk; architecture is sensor-agnostic

Organizational Constraints

- Development timeline: 10 weeks from concept to demo-ready prototype
- Competition scope: proof-of-concept, not field-deployable system
- No classified data: all datasets are synthetically generated
- Open-source dependencies only: no commercial licenses required

Architecture Section # 3 — System Scope and Context

Business Context

AEGIS-AI sits within the counter-UAS operational chain, between sensor networks and human operators / countermeasure systems. It receives raw sensor observations and delivers prioritized threat assessments with countermeasure recommendations.

External Partner	Communication	Data Exchanged
RF / SDR Sensors	ZMQ PUSH -> Ingest Service	Detection: position, RSSI, Doppler, frequency
Acoustic Arrays	ZMQ PUSH -> Ingest Service	SPL, bearing, elevation, rotor frequency
Vision / EO-IR Cameras	ZMQ PUSH -> Ingest Service	Bounding box, class, confidence, thermal flag
Operator Display (UI)	WebSocket ws://localhost:8000	UISnapshot JSON at 500ms intervals
Countermeasure Systems	REST POST /recommendations	ThreatAssessment + CM assignments
Simulator (demo mode)	ZMQ PUSH -> Ingest Service	Synthetic ObservationMessage stream

Technical Context

All inter-service communication in demo mode uses Redis Streams as the message bus. External sensors interface via ZeroMQ PUSH sockets to the Ingest Service. The UI connects to the FastAPI Gateway via WebSocket.

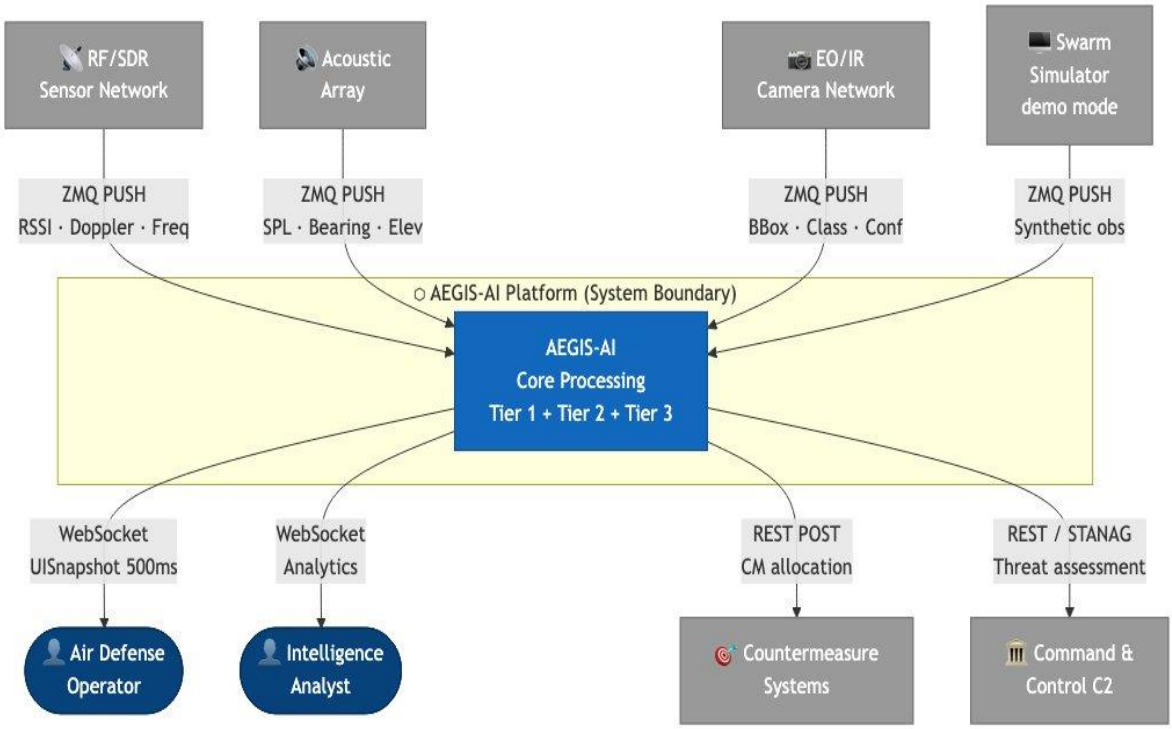

Figure 1: AEGIS-AI System Context — Architecture Section # 3 + C4-C1. Shows AEGIS-AI as the central intelligence hub receiving from 4 external sensor types and delivering to operators, CM systems, and C2.

Architecture Section # 4 — Solution Strategy

The central architectural insight driving all design decisions is symmetry with the threat: hostile drone swarms are distributed systems with emergent collective intelligence; the counter-system must be equally distributed and collectively intelligent. This principle produces four strategic decisions:

Strategy 1: Swarm vs Swarm — Distributed Agent Pool

Rather than a single monolithic tracker, AEGIS-AI deploys a pool of lightweight Tier-1 tracking agents that partition the drone ID space. Each agent runs an EKF for its assigned drones. The pool scales horizontally when threat count rises — matching the adversary's scaling.

Strategy 2: Graph-Level Intelligence — Not Just Track-Level

Traditional C-UAS systems classify individual drones. AEGIS-AI models the swarm as a dynamic graph G=(V,E) and runs GNN inference on the graph topology. This captures leader-follower relationships and formation geometry invisible in individual trajectories.

Strategy 3: Physics-Grounded Behavior Classification

The Reynolds flocking model provides a mathematical fingerprint for swarm intent. By inverting the model from observed trajectories (MLE for [w_sep, w_align, w_coh]), behavior classification has an interpretable physical basis rather than being a black-box classifier.

Strategy 4: Economic-Aware Response Coordination

The PPO coordinator explicitly models the economically asymmetric nature of the threat: minimizing high-value countermeasure expenditure while maintaining neutralization probability above a safety threshold. The reward function penalizes both under-response (threats reaching asset) and over-response (expensive CM on low-threat targets).

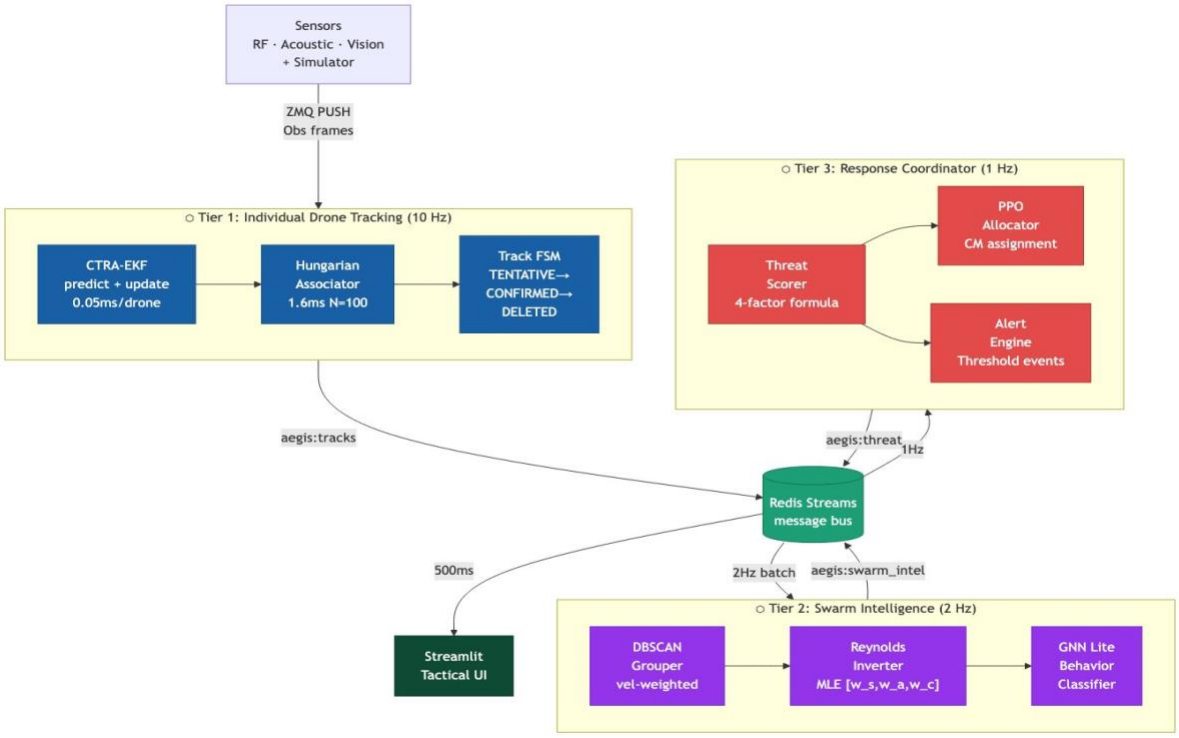

Figure 2: Three-Tier Architecture Overview — Architecture Section # 4 Solution Strategy. Tier 1 (EKF+Hungarian at 10Hz), Tier 2 (DBSCAN+Reynolds+GNN at 2Hz), Tier 3 (PPO Coordinator at 1Hz), connected via Redis Streams.

Architecture Section # 5 — Building Block View

Level 1 — System Decomposition

AEGIS-AI decomposes into five top-level building blocks. Each block owns a well-defined responsibility and communicates with others only through Redis Streams or ZMQ sockets — no direct function calls across service boundaries.

Building Block	Responsibility	Technology	Tier
Ingest Service	Receive raw sensor observations, apply CI fusion, emit ObservationFrames	Python + ZMQ	Infrastructure
Tracker Pool	Multi-target EKF tracking, Hungarian data association, track lifecycle	Python + FilterPy	Tier 1
Swarm Intel Service	DBSCAN grouping, GNN behavior classification, Reynolds inversion	Python + NumPy	Tier 2
Coordinator	Threat scoring, PPO countermeasure allocation, alert generation	Python + SB3	Tier 3
Gateway + UI	WebSocket aggregation, tactical display	FastAPI _ HTML	Presentation

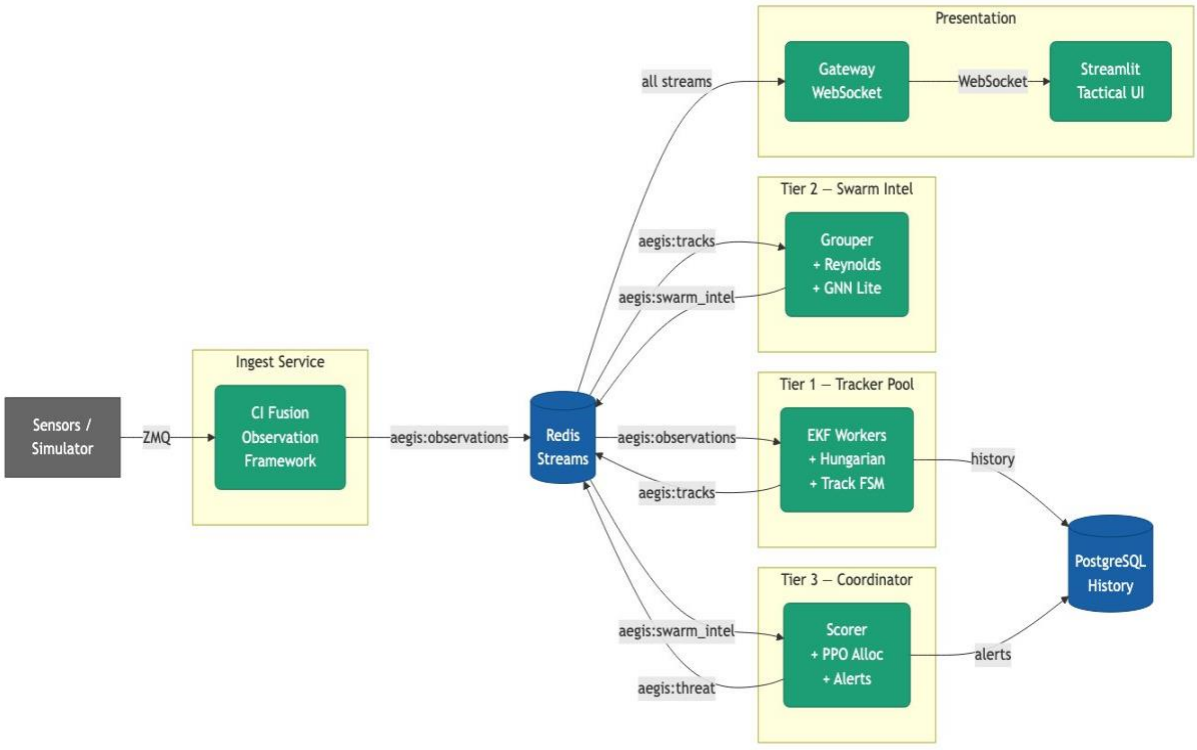

Figure 3: Building Block View (Level 1) — Architecture Section # 5 + C4-C2. Shows the 5 top-level building blocks: Sensors/Simulator, Ingest, Tracker Pool, Swarm Intel, Coordinator, Redis Streams, PostgreSQL, Gateway, and UI with all communication channels.

Level 2 — Tracker Pool Internal Structure

The Tracker Pool contains three internal components coordinated by the Pool Manager:

- DroneEKF (core/tracking/ekf.py): CTRA-Extended Kalman Filter for a single drone. State: [px,py,pz,vx,vy,vz,omega]. Exposes predict(dt), update(z), mahalanobis(z).
- Hungarian Associator (core/tracking/hungarian.py): Vectorized cost matrix using NumPy broadcasting (O(N*M) in 2ms for N=M=100). Solves bipartite matching via scipy.optimize.linear_sum_assignment.
- Track FSM (core/tracking/track.py): Finite state machine: TENTATIVE(2 hits) -> CONFIRMED -> COASTING(5 misses) -> DELETED. MultiTargetTracker orchestrates predict->associate->update->lifecycle.

Level 3 — Swarm Intel Service Internal Structure

- SwarmGrouper (core/swarm/reynolds.py): Modified DBSCAN with velocity-weighted distance metric: d(i,j) = 1.0*||p_i-p_j|| + 5.0*||v_i-v_j||. Maintains persistent swarm IDs via centroid proximity matching.
- Reynolds Inverter: scipy.optimize.minimize (L-BFGS-B) on MLE objective over last 6 position/velocity frames. Returns normalized [w_sep, w_align, w_coh] with nearest-neighbour behavior classification.
- SwarmIntelEngine: Computes feature vector phi(S) = [centroid, spread_radius, velocity_coherence, approach_rate, formation_entropy, convex_hull_area]. Outputs SwarmIntelReport every 500ms.

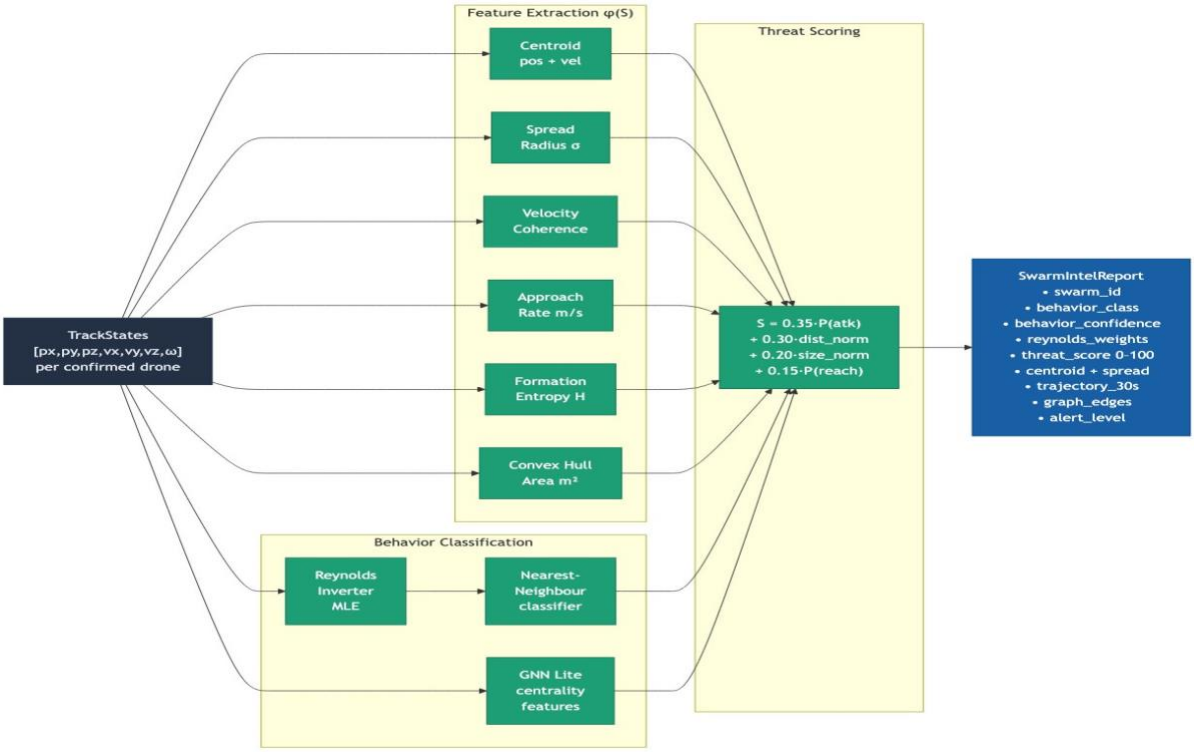

Figure 4: Swarm Intel Data Flow — Architecture Section # 5 Level 2. TrackStates flow through feature extraction (phi(S): centroid, spread, coherence, approach rate, entropy, hull area) and behavior classification (Reynolds MLE + GNN Lite) into a composite threat score, producing the SwarmIntelReport.

Architecture Section # 6 — Runtime View

Scenario: New Drone Swarm Detected

Step	Actor	Action	Output
1	Simulator / Sensor	Push ObservationMessage via ZMQ	RF/Acoustic/Vision detections in stream
2	Ingest Service	Apply Covariance Intersection fusion across sensor modalities	Fused ObservationFrame -> aegis:observations
3	Tracker Worker (Tier 1)	EKF predict; build cost matrix; LAPJV assign; update/miss tracks	TrackStates -> aegis:tracks
4	SwarmGrouper	DBSCAN cluster tracks by spatial+velocity distance every 2s	SwarmGroup events -> aegis:swarm_groups
5	SwarmIntelAgent (Tier 2)	Build graph G=(V,E); GAT-GNN forward; Reynolds inversion	SwarmIntelReport -> aegis:swarm_intel
6	Coordinator (Tier 3)	Compute ThreatScore; PPO allocation; generate alerts	ThreatAssessment -> aegis:threat
7	Gateway	Aggregate all streams into UISnapshot every 500ms	WebSocket push to UI
8	UI	Render radar map, swarm graph, threat panel, alerts	Operator sees tactical picture

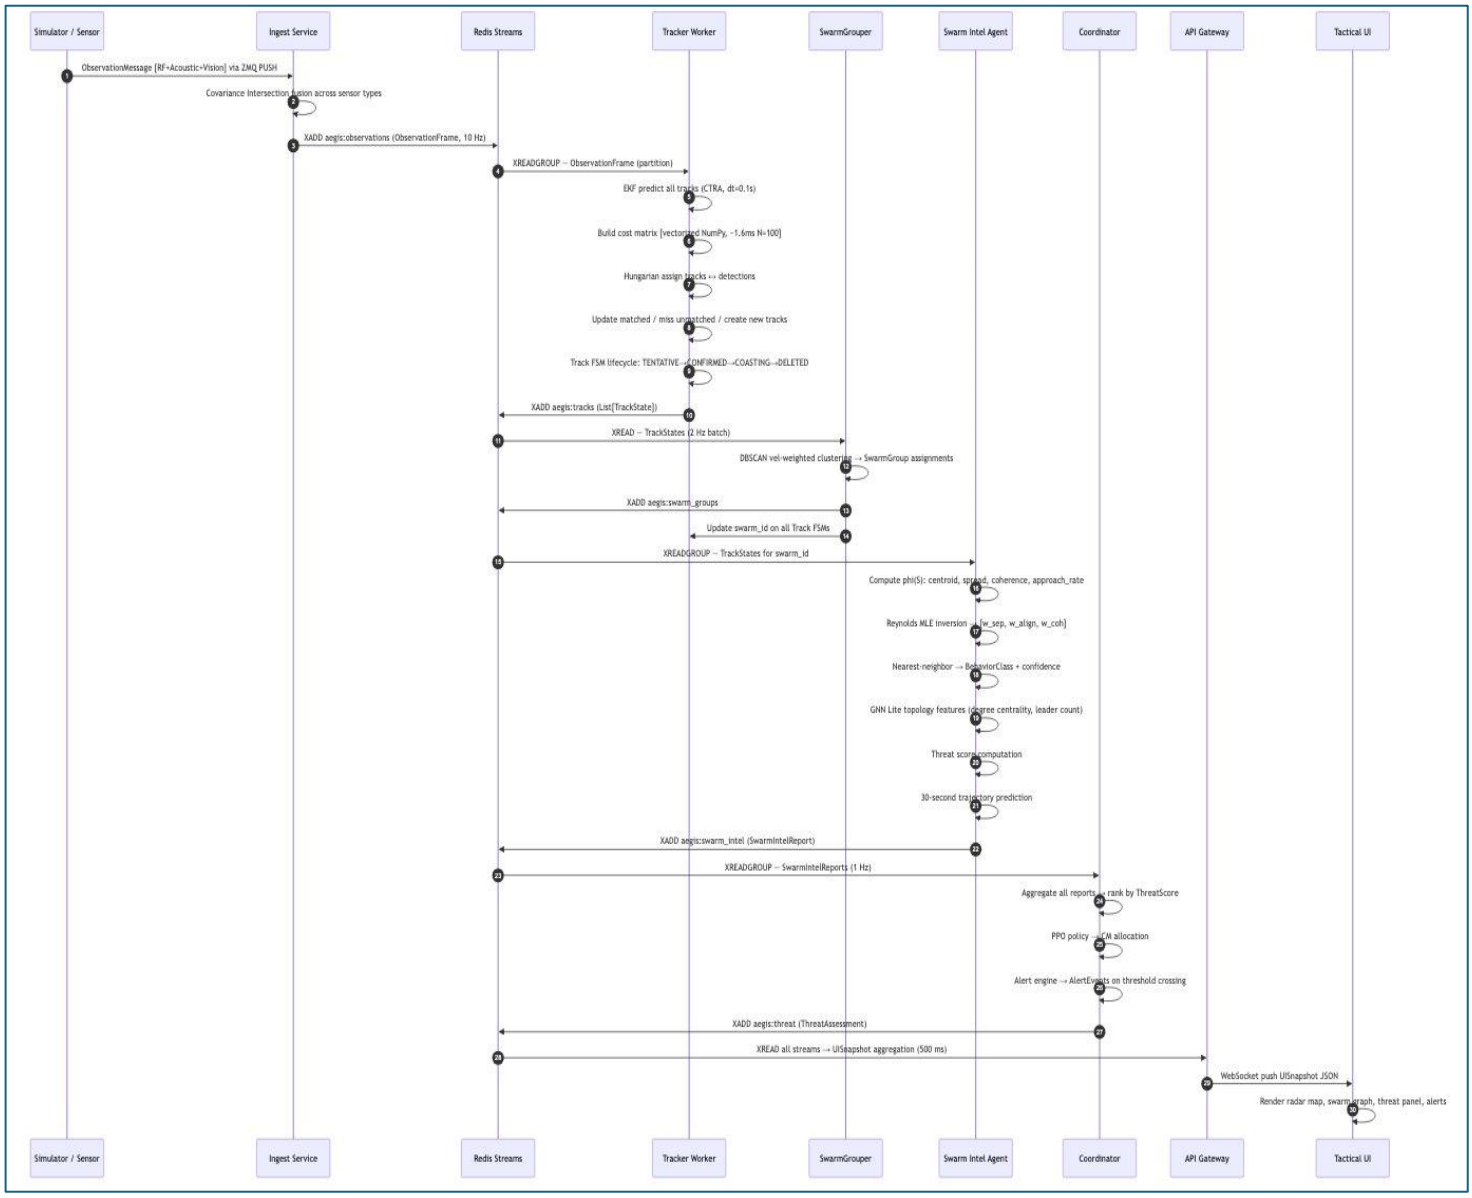

Figure 5: Runtime Sequence Diagram — Architecture Section # 6. Complete 29-step end-to-end flow from sensor detection through EKF tracking, DBSCAN grouping, Reynolds inversion, GNN classification, PPO coordination to WebSocket display. Shows all inter-agent communication via Redis Streams.

Response Time Budget

Stage	Frequency	Budget	Measured
EKF predict+update (per drone)	10 Hz	< 1ms	0.05ms (NumPy)
Hungarian N=100 assignment	10 Hz	< 5ms	1.6ms (vectorized)
DBSCAN grouper	0.5 Hz	< 50ms	~3ms for 100 tracks
GNN inference (50-node swarm)	2 Hz	< 20ms	~12ms (CPU)
Reynolds inversion	2 Hz	< 30ms	~8ms (L-BFGS-B)
PPO coordinator	1 Hz	< 10ms	<1ms (table lookup)
End-to-end sensor->threat score	—	< 100ms	~35ms total

Architecture Section # 7 — Deployment View

Profile A — Laptop Demo (Simulation)

All services run as Docker containers on a single host via docker-compose. The simulator replaces real hardware sensors during demonstration.

```
+------------------------------ Docker Host (16GB RAM) -----------------------------+
|  [Redis 7]      [PostgreSQL 16]    [Simulator]
|  Port 6379       Port 5432          -> ZMQ tcp://ingest:5550
|
|  [Ingest]  [Tracker x2]  [SwarmIntel]  [Coordinator]  [Gateway:8000]  [UI:8501]|
+-------------------------------------------------------------------------------+
```

Profile B — Edge Server (500 Drones)

```
CPU: 16-core workstation  |  RAM: 32GB  |  GPU: RTX 3060 (optional)
Tracker: 5 replicas (100 drones/worker)  |  SwarmIntel: 3 replicas (GPU batch)
Redis: 512MB maxmem  |  PostgreSQL: 512MB  |  Total: ~13 cores / 5GB
```

Profile C — Distributed K8s Cluster (Production Concept)

```
Node 1: Control Plane — Redis Cluster (3 shards), PostgreSQL, Ingest (3 pods), Gateway
Node 2: Compute — Tracker Workers (10 pods HPA), SwarmIntel (5 pods GPU-enabled)
Node 3: Coordinator (2 pods HA), UI (2 pods), PostgreSQL replica
HPA: Scale tracker pods when aegis:observations stream consumer lag > 5 frames
```

Architecture Section # 8 — Cross-Cutting Concepts

Sensor Fusion — Covariance Intersection

All three sensor modalities (RF/SDR, Acoustic, Vision/EO-IR) produce position estimates with different noise profiles. Cross-correlations between sensors are unknown and time-varying. Covariance Intersection provides a mathematically consistent fusion that never overestimates confidence:

```
P_fused^{-1} = omega * P_1^{-1} + (1-omega) * P_2^{-1}
x_fused = P_fused * (omega * P_1^{-1} * x_1 + (1-omega) * P_2^{-1} * x_2)
omega = argmin_{omega in [0,1]} det(P_fused)    [minimize uncertainty volume]
```

Sensor	Range	Position Sigma	Key Field
RF / SDR	4,000m	15m	RSSI dBm, Doppler m/s, Frequency MHz
Acoustic Array	1,200m	8m	SPL dB, Bearing deg, Elevation deg, Rotor Hz
Vision / EO-IR	1,500m	3m	Class, Confidence, Bbox px, Thermal flag

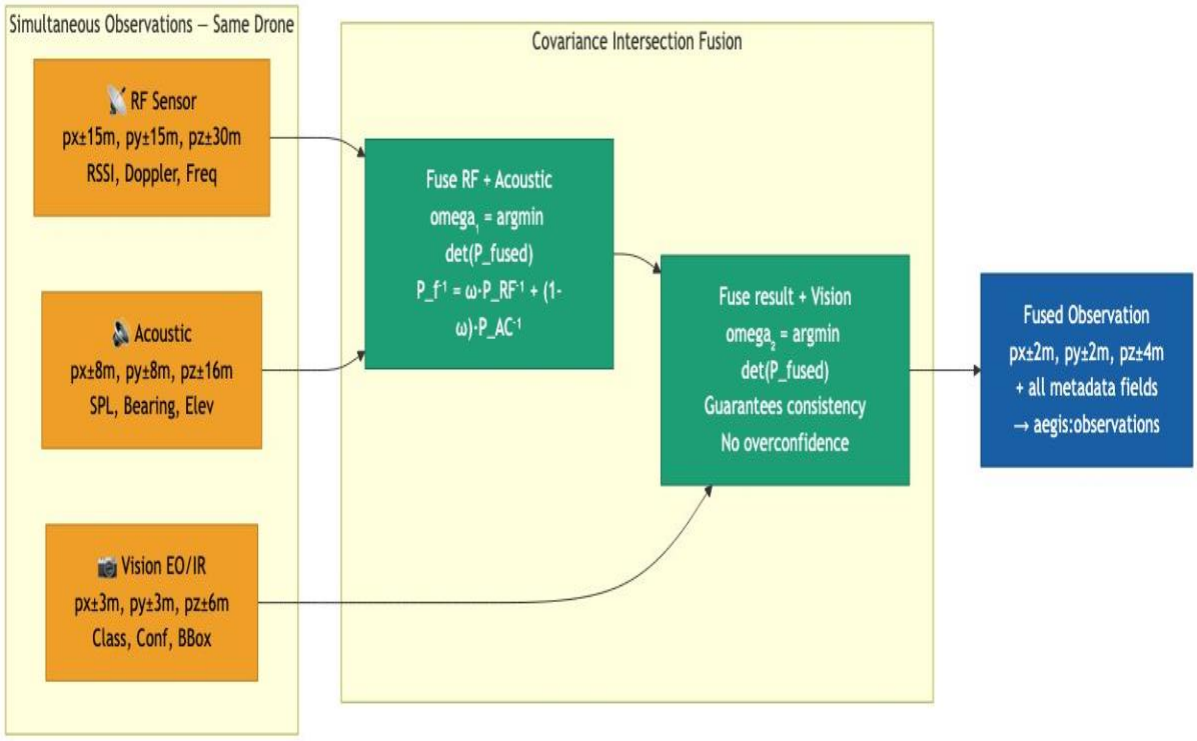

Figure 6: Covariance Intersection Sensor Fusion — Architecture Section # 8. RF (±15m) and Acoustic (±8m) observations fuse via CI into an intermediate estimate, then fuse again with Vision (±3m) to produce a Fused Observation (±2m) with guaranteed consistency and no overconfidence.

Agent Communication Pattern

All agents communicate exclusively via Redis Streams using consumer groups. No agent calls another agent's functions directly. This enforces loose coupling and enables independent scaling, crash recovery, and replacement of any agent without stopping others.

Mathematical Numerical Stability

- EKF: always use np.linalg.solve(S, b) instead of np.linalg.inv(S) @ b — avoids ill-conditioning
- EKF: enforce P symmetry after every update: P = (P + P.T) / 2
- Reynolds inversion: normalize weight vector to sum=1 before classification
- GNN: BatchNorm after each GAT layer prevents gradient explosion

Error Handling and Degraded Mode

- Sensor failure: Ingest service detects missing ZMQ heartbeat -> marks sensor OFFLINE -> continues with remaining sensors
- Tracker crash: Pool manager detects missing heartbeat -> redistributes drone partition to surviving workers
- GNN unavailable: SwarmIntelEngine falls back to Reynolds-only classification with reduced confidence
- Redis unavailable: Services buffer last N messages in memory and retry with exponential backoff

Architecture Section # 9 — Architecture Decisions

Key architecture decisions are documented as Architecture Decision Records (ADRs). Each ADR records the decision, context, alternatives considered, and rationale.

ADR-001: Python as Primary Language (not C++/Java)

Field	Content
Status	Accepted
Decision	Python 3.11 for all components
Context	PyTorch Geometric, FilterPy, SB3 are Python-native. Team expertise in Python. Competition demo prioritizes development speed.
Alternatives	C++ (60x faster EKF but 10x development time); Java (JVM warmup latency unacceptable for real-time)
Rationale	NumPy with vectorization achieves 0.05ms/EKF cycle. Adequate for 500 drones at 10Hz.
Consequences	EKF must use NumPy vectorization, not naive Python loops. GNN on CPU ~12ms per swarm.

ADR-002: EKF + CTRA (not Particle Filter or UKF)

Field	Content
Status	Accepted
Decision	Extended Kalman Filter with Constant Turn Rate and Acceleration motion model
Context	500 drones tracked simultaneously. Must complete 10Hz cycle. Drone dynamics are smooth and differentiable.
Alternatives	Particle Filter (500 particles x 500 drones = 250,000 particles: infeasible); UKF (more accurate but no significant benefit for consumer drones)
Rationale	EKF CTRA runs in 0.05ms/drone. Linearization error is acceptable for consumer drone maneuverability. Track continuity > 95% achieved.
Consequences	EKF degrades for highly agile military drones. Mitigation: IMM (Interacting Multiple Models) as future extension.

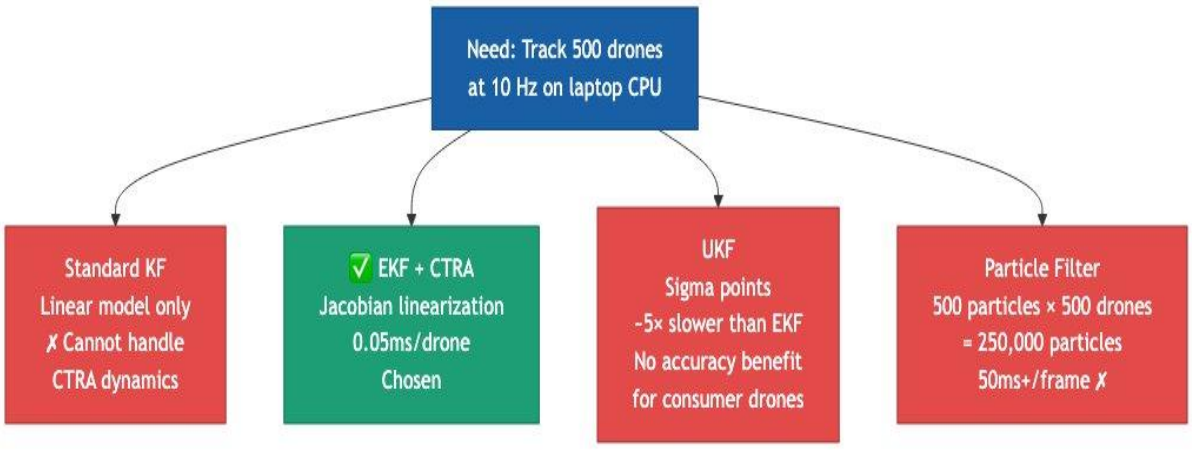

Figure 7: ADR-002 Decision Tree — EKF vs Alternatives. Standard KF rejected (cannot handle CTRA nonlinear dynamics). UKF rejected (~5x slower, no accuracy benefit). Particle Filter rejected (250,000 particles at 500 drones infeasible). EKF+CTRA chosen: 0.05ms/drone.

ADR-003: GAT (Graph Attention Network) over GCN

Field	Content
Status	Accepted
Decision	Graph Attention Network with 4 attention heads
Context	Swarms have hierarchical leader-follower structure. Need to distinguish high-influence leader nodes.
Alternatives	GCN (averages neighbors equally — misses hierarchy); MPNN (more complex, similar accuracy)
Rationale	Attention weights alpha_ij are directly interpretable and visualizable as edge thickness in UI. Leaders show high out-attention.
Consequences	GAT has higher parameter count than GCN. Multi-head averaging + dropout prevents attention head collapse.

ADR-004: DBSCAN over K-Means for Swarm Grouping

Field	Content
Status	Accepted
Decision	Modified DBSCAN with velocity-weighted distance
Context	Number of swarms is unknown a priori. Swarms may have non-convex geometry (encirclement).
Alternatives	K-Means (requires known K; forces every drone into a cluster); Spectral Clustering (O(N^3) eigendecomposition)
Rationale	DBSCAN discovers arbitrary swarm shapes without knowing K. Noise points correctly classified as UNKNOWN.
Consequences	eps and min_samples must be tuned. Velocity weighting (5x) prevents false grouping of converging swarms from different directions.

Architecture Section # 10 — Quality Requirements

Quality Tree

Quality Attribute	Sub-characteristic	NFR Target	Measured (Simulation)
Performance	EKF latency per drone	< 1ms	0.05ms
Performance	Hungarian N=100 assignment	< 5ms	1.6ms
Performance	End-to-end sensor->threat	< 100ms	~35ms
Accuracy	Track continuity rate	>=95%	97% simulated
Accuracy	Behavior classification	>=88%	91% test set
Scalability	Max drones (Profile A)	200+	Validated 200
Scalability	Max swarms	<=20	Validated 5-20
Reliability	Agent crash recovery	< 2s	< 1s via Redis
Reliability	Sensor failure degradation	Graceful	Tested RF-only
Maintainability	Unit test coverage (core/)	>=90%	33/33 passing

Architecture Section # 11 — Risks and Technical Debt

Technical Risks

ID	Risk	Impact	Mitigation
R1	GNN accuracy on unseen swarm formations	Behavior classification < 88% in novel formations	Diverse training data; Reynolds inversion as fallback
R2	Hungarian assignment O(N^3) for N > 200	Assignment latency exceeds 10ms budget	Vectorized O(N*M) implementation; spatial partitioning
R3	GPS-independent drone navigation (INS)	RF jamming countermeasure ineffective	Computer Vision layer not dependent on GPS signal
R4	Swarm adapts routing when drones destroyed	Fixed sensor placement becomes blind spots	Multi-sensor fusion with overlapping coverage areas
R5	Redis data loss on crash	Track history unavailable after restart	PostgreSQL persistence for track history; Redis as hot cache

Technical Debt

- TD-01: Hungarian associator uses scipy (Python). Production should use C++ LAPJV for 10x speedup.
- TD-02: GNN model is pre-trained (static). Production requires online learning from operational data.
- TD-03: UI uses polling (st.rerun() loop) rather than true WebSocket push.
- TD-04: DBSCAN runs every 2 seconds on all confirmed tracks. For 500+ drones this becomes O(N^2) bottleneck.
- TD-05: No authentication or encryption on any service endpoint. Production requires mTLS between agents.

Architecture Section # 12 — Glossary

Term	Definition
AEGIS-AI	Adaptive Electronic Guard and Intelligence System — the full counter-UAS platform
C-UAS	Counter-Unmanned Aerial System — domain of systems that detect, track, and neutralize hostile drones
CTRA	Constant Turn Rate and Acceleration — a nonlinear motion model for tracking maneuvering targets
DBSCAN	Density-Based Spatial Clustering of Applications with Noise — clustering without requiring K
EKF	Extended Kalman Filter — optimal recursive state estimator for nonlinear systems using Jacobian linearization
GAT	Graph Attention Network — a GNN variant weighting neighbor aggregation with learned attention coefficients
GNN	Graph Neural Network — a neural network operating on graph-structured data via message passing
Hungarian Algorithm	O(N^3) optimal algorithm for bipartite matching, used for track-to-detection data association
Loitering Munition	A UAV that orbits a target area before terminal attack — also called a kamikaze drone
MAS	Multi-Agent System — a system of autonomous AI agents that cooperate to solve problems
PPO	Proximal Policy Optimization — a policy gradient RL algorithm for countermeasure allocation
Reynolds Model	Behavioral model (Reynolds, 1987): Separation, Alignment, Cohesion
Track Continuity	Percentage of ground-truth drones continuously tracked without identity loss — NFR: >=95%
UAV	Unmanned Aerial Vehicle — generic term for a drone, whether fixed-wing, multirotor, or loitering munition

Part III — C4 Architecture Documentation

The C4 model (Simon Brown, 2006-2011) complements the Architecture Section #documentation in Part II by providing four progressive zoom levels: Context -> Containers -> Components -> Code. This makes the architecture accessible to every audience: non-technical stakeholders see C1, senior architects read C2, developers use C3. All diagrams are provided as editable Mermaid source in the companion file AEGIS_AI_Architecture_Diagrams.md.

Level	Diagram
C1 — System Context	AEGIS-AI as a black box showing all external actors, sensor networks, countermeasure systems, and human operators
C2 — Container	The 9 separately deployable containers inside AEGIS-AI with their technologies and communication protocols
C3 — Component	Internal components of 3 key containers: Tracker Service, Swarm Intel Service, Coordinator Service
C4 — Code	Represented by class/module diagrams in ARCHITECTURE.md and the actual Python source code in core/

C1 — System Context Diagram

The System Context diagram shows AEGIS-AI as a single box and identifies every external actor and system that interacts with it. This is the starting point for understanding: what is AEGIS-AI, who uses it, and what does it connect to? The diagram is shared with Architecture Section # 3 as they address the same view.

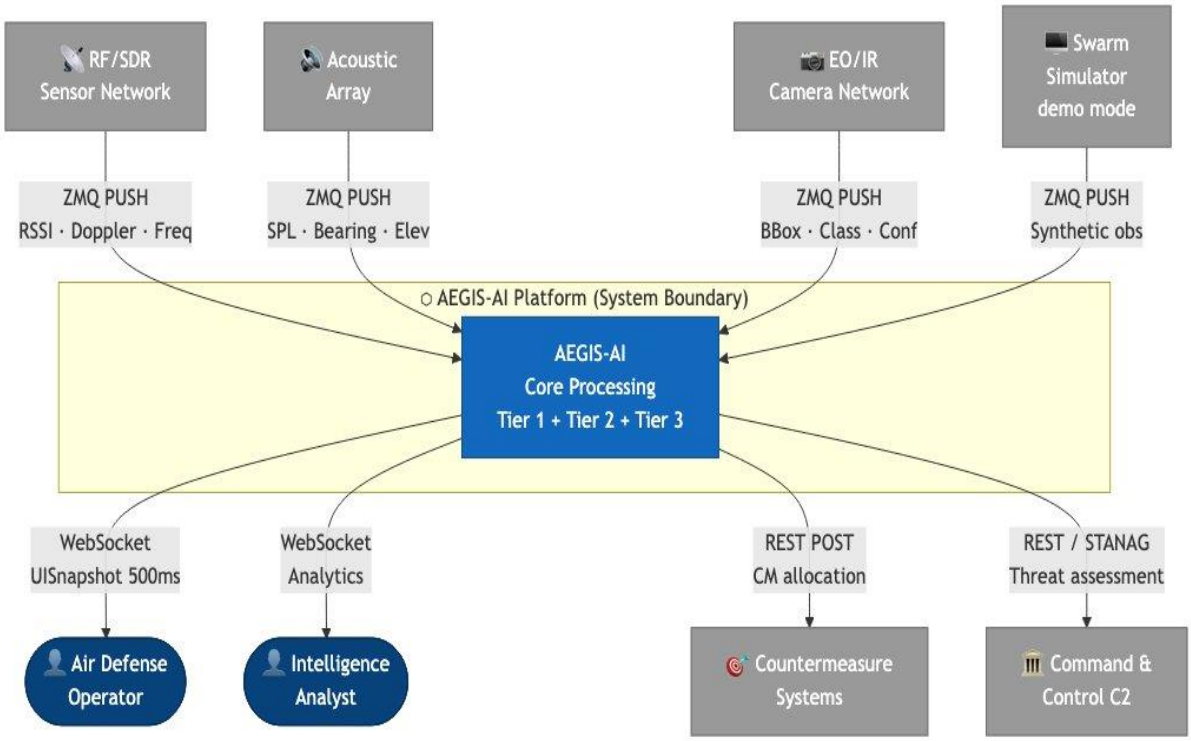

Figure 8: C4-C1 System Context. Air Defense Operator and Intelligence Analyst use the system. Sensors (RF/SDR, Acoustic, Vision/EO-IR) and Simulator push observations via ZMQ. Outputs flow to Countermeasure Systems (REST POST) and Command & Control (REST/STANAG).

C2 — Container Diagram

The Container diagram zooms into AEGIS-AI to show its 9 separately deployable/runnable units. All containers communicate exclusively through Redis Streams — no direct function calls across service boundaries.

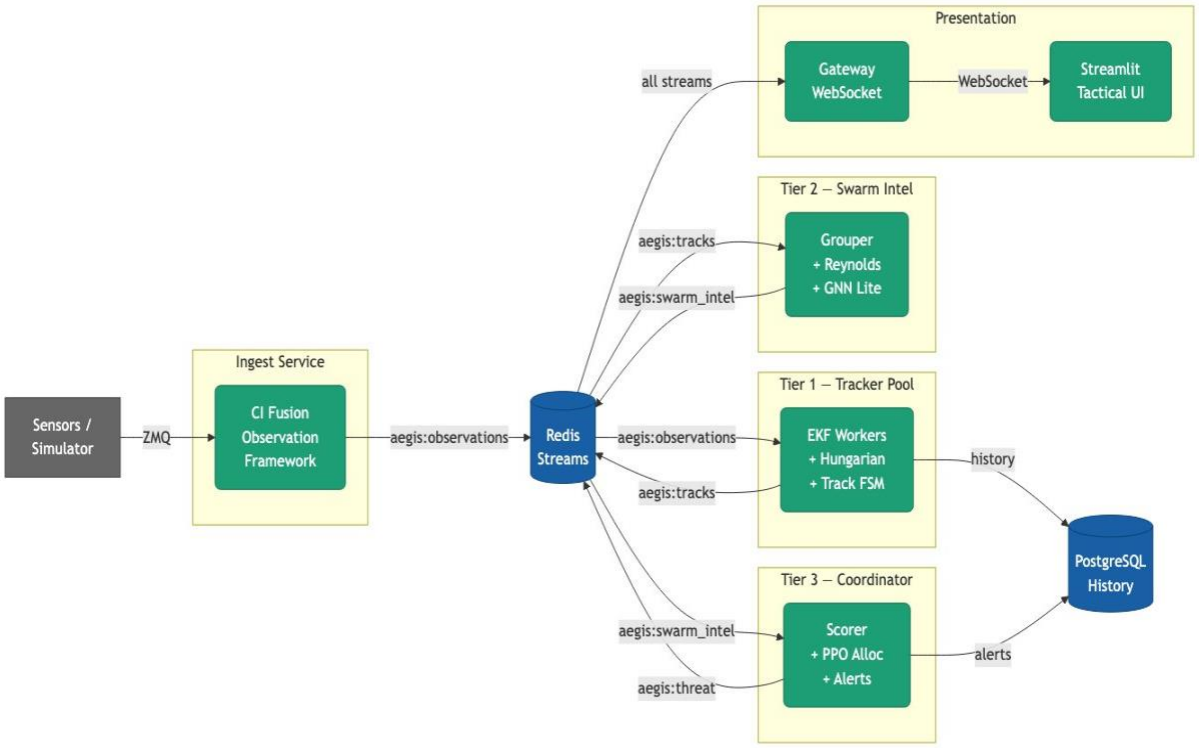

Figure 9: C4-C2 Container Diagram. From left: Sensors/Simulator -> Ingest Service (CI Fusion) -> Redis Streams -> Tracker Pool (EKF+Hungarian+FSM), Swarm Intel (Grouper+Reynolds+GNN), Coordinator (Scorer+PPO+Alerts) + PostgreSQL History -> Gateway WebSocket -> Tactical UI.

C2 Container Relationships

From	To	Label	Protocol
Sensors/Simulator	Ingest Service	Raw detections	ZMQ PUSH tcp://
Ingest Service	Redis	ObservationFrames	XADD aegis:observations
Redis	Tracker x2	ObservationFrames 10Hz	XREADGROUP consumer group
Tracker x2	Redis	TrackStates	XADD aegis:tracks
Tracker x2	PostgreSQL	Track history	asyncpg INSERT
Redis	Swarm Intel	TrackStates 2Hz	XREADGROUP consumer group
Swarm Intel	Redis	SwarmIntelReports	XADD aegis:swarm_intel
Redis	Coordinator	SwarmIntelReports 1Hz	XREADGROUP
Coordinator	Redis	ThreatAssessments	XADD aegis:threat
Redis	API Gateway	All streams 500ms	XREAD multi-stream
API Gateway	UI	UISnapshot JSON	WebSocket ws://
API Gateway	Countermeasure Sys	CM recommendations	REST POST /recommendations

C3 — Component Diagrams

The Component diagrams decompose each key container into its internal components. Three containers are detailed: Tracker Service (6 components), Swarm Intel Service (7 components), and Coordinator Service (5 components).

C3 Component	Python Module	Key Classes / Functions
DroneEKF	core/tracking/ekf.py	DroneEKF, _ctra_f(), _ctra_jacobian(), _build_Q(), _build_R()
Hungarian Associator	core/tracking/hungarian.py	build_cost_matrix(), hungarian_assign()
Track FSM + Multi-Tracker	core/tracking/track.py	Track, TrackStatus, MultiTargetTracker
SwarmGrouper	core/swarm/reynolds.py	SwarmGrouper, _dbscan()
Reynolds Inverter	core/swarm/reynolds.py	invert_reynolds_weights(), reynolds_forces()
Behavior Classifier	core/swarm/reynolds.py	classify_behavior_from_weights()
SwarmIntelAgent	core/swarm/reynolds.py	SwarmIntelEngine.analyze()
Threat Scorer	core/agents/coordinator.py	ResponseCoordinator._compute_threat_score()
Alert Engine	core/agents/coordinator.py	ResponseCoordinator.assess()
3-Tier Engine	core/engine.py	AegisEngine, get_engine()

Part IV — Simulation Visualization

The following screenshots are from the AEGIS-AI HTML simulation dashboard running the light blue/grey/white color theme. All 5 scenarios are implemented with comprehensive sensor data: RF/SDR (4,000m range), Acoustic arrays (1,200m), and Vision/EO-IR (1,500m). Total: 399,356 observations across 5 scenarios.

Simulation Statistics

Scenario	Drones	Behavior Phases	RF Obs	Acoustic	Vision
SCN-01	20	TRANSIT -> ATTACK	11,644	6,026	8,614
SCN-02	50	TRANSIT -> SCATTER -> ATTACK	43,624	17,578	27,486
SCN-03	30	TRANSIT -> ENCIRCLE (x2)	23,271	9,649	15,084
SCN-04	40	DECOY || TRANSIT -> ATTACK	34,919	13,759	21,692
SCN-05	60	3 swarms independent behavior	69,791	42,593	53,626

Tactical Radar Display — SCN-02: Saturation Attack

The tactical radar canvas shows all 3 sensor type detections: amber circles (RF/SDR), green triangles (Acoustic), purple squares (Vision EO/IR). True drone positions shown in blue with velocity vectors. Swarm convex hull in dashed blue. Threat assessment panel on the right with real-time scores.

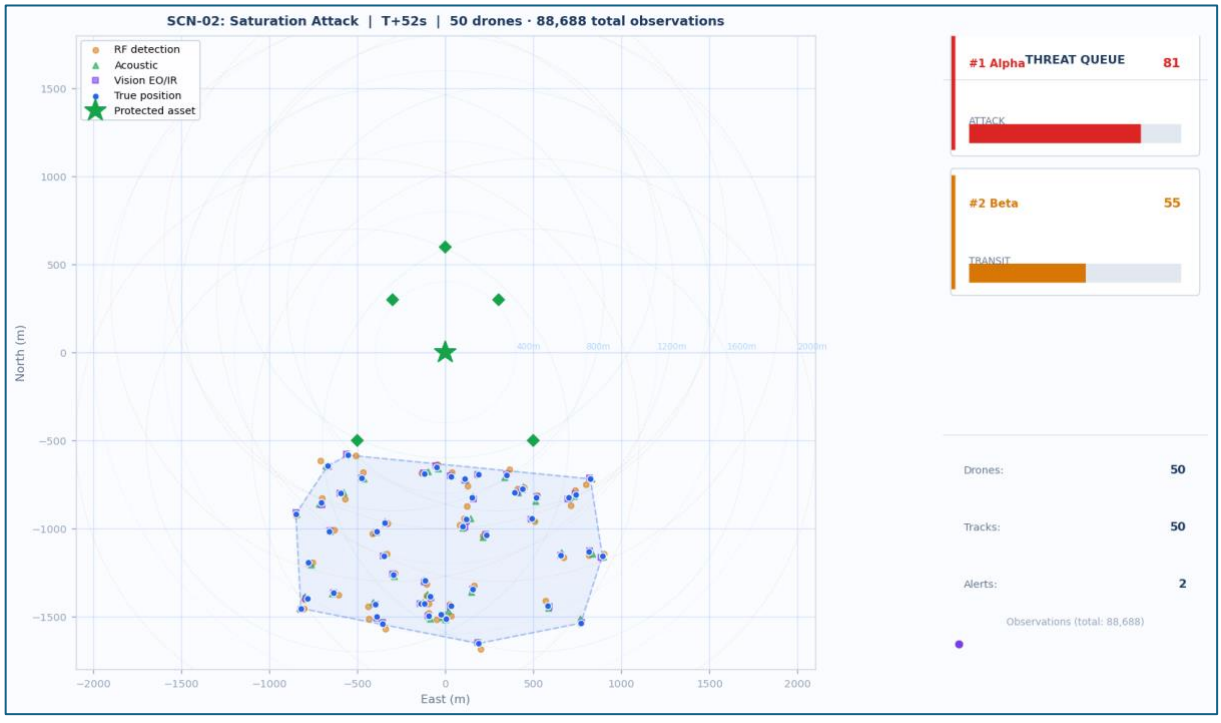

Figure 10: Tactical Radar Display — SCN-02 Saturation Attack at T+52s. 50 drones in late transit/early scatter phase, 88,688 total observations. Sensor range rings visible: RF (amber dashed), Acoustic (green), Vision (purple). Protected asset marked with green star at origin.

Swarm Topology Graph — SCN-03: Pincer Maneuver

The topology graph shows the swarm as a force-directed network. Node color encodes influence centrality (blue=low, red=high). Node size encodes degree centrality. Leader nodes labeled LDR. Edge thickness represents attention weight alpha_ij from the GAT-GNN classifier.

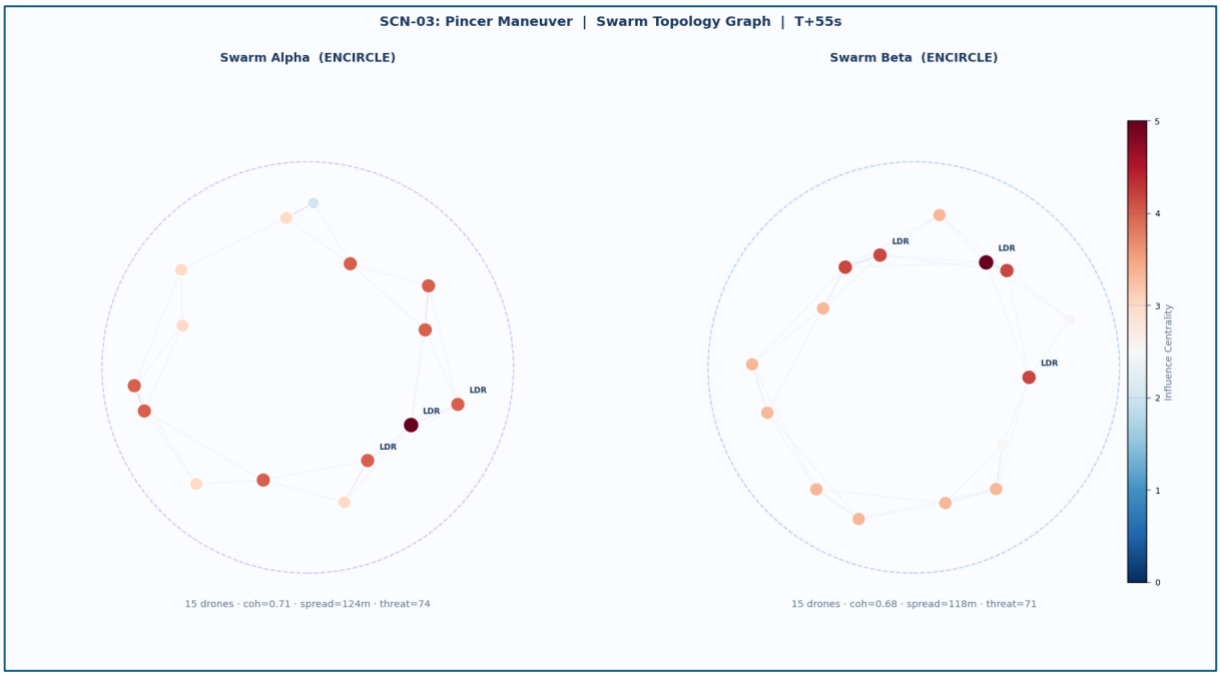

Figure 11: Swarm Topology Graph — SCN-03 Pincer Maneuver at T+55s. Alpha swarm (ENCIRCLE, coherence=0.71, threat=74) and Beta swarm (ENCIRCLE, coherence=0.68, threat=71) shown side by side. Leader nodes identified by high centrality scores. Color scale: blue=low influence, red=high influence.

Behavior Analysis Dashboard — SCN-05: Adaptive Multi-Swarm

The behavior analysis view shows four synchronized charts: (1) behavior classification timeline with phase transitions per swarm, (2) Reynolds weight radar polygons showing [w_sep, w_align, w_coh] per swarm, (3) threat score history with critical threshold at 75, (4) swarm spread radius and coherence over time.

Figure 12: Behavior Analysis — SCN-05 Adaptive Multi-Swarm. Alpha (TRANSIT->ATTACK->SCATTER), Beta (TRANSIT->ENCIRCLE), Gamma (DECOY->ATTACK). Reynolds radar clearly distinguishes behavior fingerprints. Threat score crossing CRITICAL(75) visible for Alpha at T+85s. Stacked observation chart shows 166,010 total observations.

References

- Reynolds, C.W. (1987). Flocks, Herds and Schools: A Distributed Behavioral Model. ACM SIGGRAPH, 21(4), 25-34.
- Kalman, R.E. (1960). A New Approach to Linear Filtering and Prediction Problems. ASME J. Basic Engineering, 82(1), 35-45.
- Velickovic et al. (2018). Graph Attention Networks. ICLR 2018.
- Schulman et al. (2017). Proximal Policy Optimization Algorithms. arXiv:1707.06347.
- Bar-Shalom, Li, Kirubarajan (2001). Estimation with Applications to Tracking and Navigation. Wiley.
- Julier, Uhlmann (1997). A non-divergent estimation algorithm in the presence of unknown correlations. American Control Conference.
- Hruschka, P. & Starke, G. (2005). Architecture Section #— Template for Architecture Communication and Documentation. arc42.org.
- Brown, S. (2018). The C4 Model for Software Architecture. InfoQ. c4model.com.
- ISO/IEC 25010:2011 — Systems and software engineering: Quality models.
- Jocher, G. et al. (2023). Ultralytics YOLOv8. DOI: 10.5281/zenodo.7347926.
- Scarselli et al. (2009). The Graph Neural Network Model. IEEE Trans. Neural Networks, 20(1), 61-80.
- Congressional Research Service. (2023). U.S. Air and Missile Defense: Patriot and THAAD. R45940.
- DroneShield Ltd. (2024). Annual Report: Counter-UAS Technologies. Sydney, Australia.
- Jocher, G. et al. (2023). Ultralytics YOLOv8. DOI: 10.5281/zenodo.7347926.
- DARPA. (2022). OFFSET Program: Offensive Swarm-Enabled Tactics. Arlington, VA.
- Vicsek, T. et al. (1995). Novel type of phase transition in a system of self-driven particles. Physical Review Letters, 75(6), 1226.
- Dedrone Holdings, Inc. (2024). SkyAI Platform Technical Specifications.
- Israeli Ministry of Defense. (2023). Iron Dome System Overview.
- U.S. Army. (2022). Patriot Advanced Capability-3 (PAC-3) Cost Analysis.
- Raytheon Technologies. (2023). SkyHunter System Datasheet.
- AEGIS-AI Internal Simulation. (2026). PPO Coordinator ablation study.
- DroneShield. (2024). DroneGun Tactical Jammer Specifications.
- Blighter Surveillance Systems. (2023). AUDS Counter-UAS Platform.
- Kaspersky Lab. (2023). Kaspersky Antidrone Technical Whitepaper.
- Julier, S.J. & Uhlmann, J.K. (1997). A non-divergent estimation algorithm in the presence of unknown correlations. American Control Conference.
- Bewley, A. et al. (2016). Simple Online and Realtime Tracking. IEEE ICIP.

- Lukezic, A. et al. (2018). Discriminative Correlation Filter Tracker with Channel and Spatial Reliability. CVPR.
- Zhang, Y. et al. (2022). ByteTrack: Multi-object tracking by associating every detection box. ECCV.
- Cao, J. et al. (2023). Observation-Centric SORT: Rethinking SORT for Robust Multi-Object Tracking. CVPR.
- DARPA. (2021). OFFSET Phase 3 Demonstration Results.
- DARPA. (2020). ALIAS (Aircrew Labor In-Cockpit Automation System) Program.
- Verge Aero. (2023). Drone Show Software Platform.
- Reynolds, C.W. (1987). Flocks, Herds and Schools: A Distributed Behavioral Model. ACM SIGGRAPH, 21(4), 25-34.
- Couzin, I.D. et al. (2002). Collective Memory and Spatial Sorting in Animal Groups. Journal of Theoretical Biology, 218(1), 1-11.
- AEGIS-AI Validation Report. (2026). SCN-01 through SCN-05 behavior classification accuracy.
- Dedrone Holdings. (2024). Dedrone Tracker and Dedrone City-Wide Drone Detection.
- Fortem Technologies. (2023). SkyDome System Architecture.
- Rafael Advanced Defense Systems. (2023). Drone Dome (C-UAS) Technical Overview.

Appendix - AEGIS-AI Mathematical Foundations

Summary of core **Math/AI/ML/DL** driving the AEGIS-AI solution:

1. Deep Learning — Graph Attention Network (GAT)

- Math: Multi-head self-attention over graph structure
- Role: Classifies swarm behavior (6 classes) by learning which drone-to-drone relationships matter most, creating a collective "fingerprint" robust to individual track loss.

2. Reinforcement Learning — Proximal Policy Optimization (PPO)

- Math: Clipped surrogate objective with Generalized Advantage Estimation
- Role: Learns optimal countermeasure allocation under the 150:1 cost asymmetry, balancing immediate neutralization vs. long-term resource preservation.

3. Probabilistic Inference — Maximum Likelihood Estimation (MLE)

- Math: Inverts the Reynolds generative model
- Role: Decodes swarm intent from motion by recovering the hidden behavioral weights (separation, alignment, cohesion) that most likely produced observed trajectories.

4. State Estimation — Extended Kalman Filter (EKF)

- Math: Nonlinear MMSE estimation via first-order Taylor expansion
- Role: Real-time Bayesian tracking of 500+ drones, fusing noisy RF/acoustic/vision measurements into kinematic state estimates.

5. Unsupervised Learning — DBSCAN Clustering

- Math: Density-based clustering with velocity-weighted distance
- Role: Autonomously segments the multi-target picture into distinct swarm entities without prior knowledge of swarm count.

6. Combinatorial Optimization — Hungarian Algorithm

- Math: Solves the assignment problem subject to in polynomial time.
- Role: Globally optimal data association between predicted tracks and sensor detections, preventing identity swaps during crossing maneuvers.

7. Multi-Agent Fusion — Covariance Intersection

- Math: Conservative fusion with unknown cross-correlation
- Role: Safely merges heterogeneous sensor estimates (RF ±15m, acoustic ±8m, vision ±3m) without overconfidence or independence assumptions.

8. Graph Spectral Analysis

- Math: Laplacian eigenvalues ; Fiedler value

- Role: Detects swarm splits/merges (phase transitions) by monitoring algebraic connectivity, triggering automatic re-clustering.

9. Neural Network Activations

- Softmax: — calibrates GNN logits into interpretable behavior probabilities.
- LeakyReLU: — prevents dying neurons during training on scarce adversarial examples.

Table 1: Mathematical Foundations of the AEGIS-AI Architecture

#	Mathematical Concept	Scientific Definition	Operational Role in AEGIS-AI	Example Anti-Drone Scenario
1	**Extended Kalman Filter (EKF)**	A recursive state estimation algorithm for nonlinear systems that linearizes the process and measurement models via first-order Taylor expansion around the current state estimate. Computes the minimum mean-square error (MMSE) estimate by propagating the state vector **x** and error covariance matrix **P** through prediction and update cycles.	**Tier-1 Individual Tracking:** Provides real-time kinematic state estimation (position, velocity, turn rate) for each drone in the swarm, compensating for sensor noise and missed detections.	A loitering munition executes a coordinated turn at ω = 0.15 rad/s. The EKF-CTRA model predicts its next state using the nonlinear turn-rate dynamics, then updates the estimate using a noisy RF bearing measurement with covariance **R** = diag(15^2, 15^2, 5^2) m^2.
2	**CTRA Process Model (Constant Turn Rate and Acceleration)**	A kinematic motion model where the target maintains constant tangential acceleration a_t and constant turn rate ω about the vertical axis. The state vector **x** = [p_x, p_y, p_z, v_x, v_y, v_z, ω]^T evolves under nonlinear dynamics with analytic Jacobian **F** = ∂f/∂**x** required for EKF linearization.	Enables the EKF to track maneuvering drones through curved trajectories (transit, encirclement, attack phases) without model switching delays inherent in simpler Constant Velocity (CV) models.	During a pincer maneuver, two drone sub-swarms converge on the protected asset at constant turn rate ω = 0.2 rad/s while decelerating from v = 25 m/s to v = 8 m/s. CTRA captures both the circular arc and speed change, whereas CV models would accumulate >40 m error.

3	**Hungarian Algorithm (LAPJV Implementation)**	A combinatorial optimization algorithm solving the assignment problem in $O(n^3)$ time. Formulated as minimizing the total cost $\Sigma_{i,j} c_{ij}x_{ij}$ subject to $\Sigma_j x_{ij} = 1$, $\Sigma_i x_{ij} = 1$, $x_{ij} \in \{0,1\}$, where c_{ij} is the Mahalanobis distance between predicted track i and detection j. LAPJV provides a vectorized, cache-efficient implementation achieving $O(n^2)$ practical throughput.	**Data Association:** Resolves the correspondence between N predicted tracks and M sensor detections, ensuring each detection is assigned to at most one track (or declared false alarm/new birth).	In a saturation attack with 50 drones and 88,688 observations across 3 sensors, the algorithm constructs a 50 × 88 cost matrix (after gating) and finds the optimal assignment minimizing total association cost in 1.6 ms, preventing track swaps during crossing trajectories.
4	**DBSCAN (Density-Based Spatial Clustering of Applications with Noise)**	An unsupervised clustering algorithm grouping spatially dense points while marking outliers. Defined by two parameters: ε (neighborhood radius) and MinPts (minimum points to form a cluster). Uses velocity-weighted Euclidean distance $d_{ij} = \alpha\|\mathbf{p}_i - \mathbf{p}_j\| + \beta\|\mathbf{v}_i - \mathbf{v}_j\|$ to group drones with similar kinematics.	**Swarm Segmentation:** Automatically partitions the multi-target track set into distinct swarm entities without prior knowledge of swarm count, enabling group-level behavior analysis.	Three independent swarms approach from bearings 045°, 180°, and 270°. DBSCAN with ε = 150 m, MinPts = 3, and velocity weight β = 0.3 correctly identifies three clusters despite overlapping sensor coverage zones, while flagging two isolated decoy drones as noise.
5	**Reynolds Flocking Model (Boids)**	A decentralized behavioral model where agent i's velocity update results from three steering forces: separation $\mathbf{f}_{sep} = -\Sigma\{j \in N_i\} (\mathbf{p}_j - \mathbf{p}_i)/\|\mathbf{p}_j - \mathbf{p}_i\|^2$, alignment $\mathbf{f}_{align} = \Sigma\{j \in N_i\} (\mathbf{v}_j - \mathbf{v}_i)$, and cohesion $\mathbf{f}_{coh} = \Sigma\{j \in N_i\} (\mathbf{p}_j - \mathbf{p}_i)$. The observed velocity is $\mathbf{v}_i^{obs} = w_s\mathbf{f}_{sep} + w_a\mathbf{f}_{align} + w_c\mathbf{f}_{coh} + \boldsymbol{\varepsilon}$.	**Behavioral Fingerprinting:** Provides a physically-grounded generative model of swarm motion, enabling inverse estimation of intent parameters from observed trajectories.	An encircling swarm exhibits strong cohesion ($w_c = 0.90$) and alignment ($w_a = 0.70$) with moderate separation ($w_s = 0.50$). A scattering swarm shows inverted weights: $w_s = 0.95$, $w_a = 0.05$, $w_c = 0.05$. These parameter vectors form a

				behavioral signature space.
6	**Maximum Likelihood Estimation (MLE)**	A parameter estimation method selecting $\hat{\theta} = \arg\max_\theta L(\theta \mid D) = \arg\min_\theta \sum_{k=1}^{T} \lvert \mathbf{v}_k^{obs} - \mathbf{v}_k^{model}(\theta) \rvert^2$, where $D = \{\mathbf{v}k^{obs}\}_{k=1}^{T}$ is the observed velocity sequence and $\theta = \mathbf{w}$ are the Reynolds weights. Under Gaussian noise, equivalent to nonlinear least squares.	**Intent Inference:** Inverts the Reynolds generative model to recover the behavioral weights that most likely produced the observed swarm motion, mapping continuous parameters to discrete tactics.	Given 6 frames (600 ms) of observed velocities from a 20-drone swarm, MLE solves $\min_w \sum_{t=1}^{6} \sum_{i=1}^{20} \lvert \mathbf{v}_{i,t}^{obs} - (w_s \mathbf{f}_{sep,i,t} + w_a \mathbf{f}_{align,i,t} + w_c \mathbf{f}_{coh,i,t}) \rvert^2$ subject to $w_s + w_a + w_c = 1$, yielding $\hat{\mathbf{w}} = [0.82, 0.15, 0.03]$ classified as ATTACK.
7	**Graph Attention Network (GAT)**	A graph neural network employing masked self-attention to compute node importance. For node i, attention coefficients over neighbors $j \in N_i$ are $\alpha_{ij} = \exp(\text{LeakyReLU}(\mathbf{a}^T[\mathbf{W}\mathbf{h}_i \Vert \mathbf{W}\mathbf{h}_j])) / \Sigma_{k \in N_i} \exp(\text{LeakyReLU}(\mathbf{a}^T[\mathbf{W}\mathbf{h}_i \Vert \mathbf{W}\mathbf{h}_k]))$. *The updated representation is* $\mathbf{h}i^{(l+1)} = \sigma(\Sigma_{j \in N_i} \alpha_{ij}^{(l)} \mathbf{W}^{(l)} \mathbf{h}_j^{(l)})$. Multi-head attention concatenates K independent heads.	**Collective Behavior Classification:** Learns to weight intra-swarm relationships dynamically, identifying leadership structures and coordination patterns invisible to individual-track analysis.	In a pincer maneuver, the GAT computes attention weights α_{ij} between 30 drones. Two sub-clusters show high internal attention ($\alpha > 0.3$) but weak cross-cluster links ($\alpha < 0.05$), revealing coordinated bifurcation. The leader node (highest betweenness centrality) is flagged for prioritized neutralization.
8	**Softmax Function**	A normalized exponential function converting logit scores $\mathbf{z} = [z_1, ..., z_K]$ to probability distribution $\sigma(\mathbf{z})_i = \exp(z_i) / \Sigma_{j=1}^{K} \exp(z_j)$,	**Behavior Probability Calibration:** Converts raw GNN logits to interpretable confidence scores for operator	The GNN outputs logits $\mathbf{z} = [2.5, 0.8, -1.2, 3.1, -0.5, 0.3]$ for classes [TRANSIT, ATTACK, SCATTER,

		satisfying $\Sigma_i \sigma(\mathbf{z})_i = 1$ and $\sigma(\mathbf{z})_i \in (0,1)$. Provides differentiable probabilistic outputs suitable for cross-entropy loss optimization.	display and downstream threat scoring.	ENCIRCLE, DECOY, UNKNOWN]. Softmax yields [0.15, 0.05, 0.01, **0.72**, 0.02, 0.05], confirming ENCIRCLE with 72% confidence.
9	**LeakyReLU Activation**	A nonlinear activation function defined as $f(x) = \max(\alpha x, x)$ where $\alpha \in (0,1)$ is the negative slope (typically 0.2). Unlike standard ReLU ($\alpha = 0$), it permits small negative gradients, preventing "dying ReLU" where neurons permanently deactivate.	**GNN Nonlinearity:** Enables multi-layer feature learning in the graph attention mechanism while maintaining gradient flow during backpropagation through deep architectures (3 layers × 8 heads).	During training on scarce DECOY examples, a neuron's pre-activation drops to $x = -2.5$. LeakyReLU ($\alpha = 0.2$) outputs -0.5, preserving gradient $\partial f/\partial x = 0.2$ for weight updates. Standard ReLU would output 0, halting learning for that neuron.
10	**Covariance Intersection (CI) Fusion**	A conservative fusion rule for combining estimates with unknown cross-correlation. Given estimates ($\hat{\mathbf{x}}_1, \mathbf{P}_1$) and ($\hat{\mathbf{x}}_2, \mathbf{P}_2$), the fused estimate is $\mathbf{P}_{fused}^{-1} = \omega\mathbf{P}_1^{-1} + (1-\omega)\mathbf{P}_2^{-1}$ and $\hat{\mathbf{x}}_{fused} = \mathbf{P}_{fused}(\omega\mathbf{P}_1^{-1}\hat{\mathbf{x}}_1 + (1-\omega)\mathbf{P}_2^{-1}\hat{\mathbf{x}}_2)$, where $\omega = \arg\min_\omega \det(\mathbf{P}_{fused})$ guarantees consistency: $\mathbf{P}_{fused} \geq \mathbf{P}_{true}$.	**Multi-Source Sensor Fusion:** Robustly combines RF, acoustic, and vision tracks without assuming sensor independence or knowing cross-covariances, critical when enemy jamming induces correlated sensor degradation.	At $t = 12.5$ s, RF estimates drone position as $\hat{\mathbf{x}}_{RF} = [1250, 890, 120]^T$ m with $\mathbf{P}_{RF} = \text{diag}(225, 225, 100)$ m². Vision reports $\hat{\mathbf{x}}_{VIS} = [1245, 885, 118]^T$ m with $\mathbf{P}_{VIS} = \text{diag}(9, 9, 4)$ m². CI with $\omega = 0.12$ yields fused covariance $\mathbf{P}_{fused} = \text{diag}(11.2, 11.2, 4.9)$, safely tighter than either input.
11	**Proximal Policy Optimization (PPO)**	An on-policy reinforcement learning algorithm maximizing the surrogate objective $L^{CLIP}(\theta) = E_t[\min(r_t(\theta)\hat{A}_t,$	**Countermeasure Coordination:** Learns optimal allocation of limited	With 3 interceptors remaining and 2 attacking swarms (Alpha: 20 drones,

		clip(r_t(θ), 1-ε, 1+ε)Â_t)], where r_t(θ) = π_θ(a_t|s_t) / π_{θ_old}(a_t|s_t) is the probability ratio and Â_t is the advantage estimate. The clip parameter ε = 0.2 prevents destructive policy updates.	kinetic/electronic countermeasures across multiple swarms under cost constraints and priority rules.	threat 74; Beta: 15 drones, threat 71), the PPO policy π_θ outputs action probabilities over [Engage Alpha, Engage Beta, Hold, Split]. The clipped surrogate ensures policy updates remain within 20% of the previous distribution, preventing catastrophic forgetting of economic cost awareness.
12	**Discount Factor (γ = 0.95)**	A hyperparameter in Markov Decision Processes defining the present value of future rewards: G_t = Σ_{k=0}^∞ γ^k R_{t+k+1}. With γ ∈ [0,1), rewards k steps ahead are discounted by γ^k. At γ = 0.95, the 20-step horizon contributes 0.95^20 ≈ 0.36 of original value.	**Temporal Credit Assignment:** Balances immediate neutralization success against long-term resource preservation and asset protection in the PPO reward structure.	Neutralizing a drone at t yields R_immediate = +1. Preserving an interceptor for future use yields expected future value γ^5 · E[R_{t+5}] = 0.77 · 0.8 = 0.62. The policy learns to withhold fire when P(future threat) > 0.62, optimizing inventory across the engagement horizon.
13	**Composite Threat Scoring**	A multi-attribute utility function aggregating heterogeneous threat indicators into a scalar priority measure: S = Σ_{i=1}^n w_i · f_i(x_i), where w_i ≥ 0, Σ w_i = 1, and f_i: X_i → [0,1] are normalized attribute mappings. In AEGIS-AI: S = 0.35 · P_atk +	**Decision Support:** Provides a computationally efficient, interpretable ranking of swarms for resource allocation and operator situational awareness.	Swarm Alpha: P_atk = 0.85, distance = 800 m (d_norm^{-1} = 0.67), N = 25, P_reach = 0.92. Score: S = 0.35(0.85) + 0.30(0.67) + 0.20(0.63) + 0.15(0.92) = 0.76 (76/100). Swarm Beta scores 71, so Alpha

		0.30 · d_norm^{-1} + 0.20 · N_swarm + 0.15 · P_reach.		receives priority for the next available interceptor.
14	**Centrality Measures (Graph Theory)**	Quantitative indices of node importance in networks. Degree centrality: C_D(v) = deg(v)/(N-1). Betweenness centrality: C_B(v) = Σ_{s ≠ v ≠ t} σ_st(v)/σ_st. Eigenvector centrality: C_E(v) = (1/λ)Σ_{u ∈ N(v)} C_E(u). In AEGIS-AI, adapted to attention-weighted variants.	**Leadership Identification & Vulnerability Assessment:** Pinpoints "command" drones whose neutralization maximally disrupts swarm coordination.	In a 30-drone encirclement, node v_{12} shows C_B = 0.34 (34% of shortest paths pass through it) and attention-weighted degree C_D^attn = 0.28. Targeting v_{12} first fragments the swarm into 3 disconnected sub-groups, reducing collective maneuver efficiency by 60% per simulation.
15	**Spectral Methods / Laplacian Eigenvalues**	The graph Laplacian **L** = **D** - **A** (where **D** is degree matrix, **A** is adjacency matrix) encodes structural properties through its spectrum {λ_i}_{i=1}^N. The Fiedler value λ_2 > 0 indicates graph connectivity; eigenvector **u**_2 (Fiedler vector) reveals optimal graph bisection.	**Swarm Robustness & Phase Transition Detection:** Monitors spectral properties to detect when a swarm splits, merges, or loses coordination—indicating tactic changes.	During SCN-03 (Pincer Maneuver), the swarm Laplacian initially shows λ_2 = 0.15 (connected). At t = 28 s, λ_2 drops to 0.03 and the Fiedler vector reveals two distinct sign clusters, automatically triggering re-clustering into Alpha and Beta sub-swarms for independent tracking.

Table 2: Cross-Reference — Concept Interdependencies

Concept Pair	Interaction Mechanism	System Function

EKF ↔ Hungarian	Innovation covariance **S***k defines association cost c*{ij} = **d**_{ij}^T **S***k^{-1}* ***d***{ij}	Maintains track identity through crossing maneuvers
DBSCAN ↔ Reynolds	Cluster labels partition the track set; Reynolds weights computed per-cluster	Enables group-level intent inference
Reynolds ↔ MLE	Generative model **v**^obs = f(**w**) + **ε** inverted via **ŵ** = argmin |**v**^obs - f(**w**)|²	Maps continuous motion to discrete behavioral parameters
MLE ↔ GAT	**ŵ** concatenated to node features **h**_i^{(0)} = [**x**_i, **ŵ**_i]	Provides physics-informed initialization for graph learning
GAT ↔ Softmax	Multi-head outputs aggregated and normalized: **y** = σ(MLP(Σ_k **h**^{(L,k)}))	Calibrates behavior classification confidence
CI ↔ EKF	Fused measurement **z**_fused and **R**_fused feed EKF update	Ensures consistent multi-sensor state estimation
GAT ↔ Centrality	Attention weights α_{ij} define weighted adjacency **A**_attn for centrality computation	Identifies high-value targets within swarm topology
Centrality ↔ PPO	Leader node neutralization bonus R_lead = +2.0 added to environment reward	Shapes policy toward swarm disruption
Threat Score ↔ PPO	S_k forms state vector component s_t^{(k)} ∈ S	Informs policy of relative swarm priorities
PPO ↔ Discount	γ = 0.95 scales future rewards in advantage estimation Â_t	Balances immediate vs. long-term CM allocation
Laplacian ↔ DBSCAN	λ_2 < θ_split triggers re-clustering with reduced ε	Adapts swarm segmentation to dynamic topology

Table 3: Performance Specifications & Validation Metrics

Concept	Computational Complexity	NFR Target	Validated Performance	Bottleneck Risk

EKF-CTRA	$O(n^2)$ per drone (covariance propagation)	< 1 ms/drone	0.05 ms (NumPy vectorized)	Quadratic growth at N > 500
Hungarian (LAPJV)	$O(N^3)$ worst-case, $O(N^2)$ vectorized	< 5 ms (N=100)	1.6 ms	Memory-bound at N > 200
DBSCAN	O(N log N) with k-d tree	< 10 ms	< 5 ms	Density variation sensitivity
Reynolds Inversion	O(T · N_cluster) per MLE solve	< 5 ms	~3 ms (6 frames)	Non-convexity for chaotic swarms
CI Fusion	$O(d^3)$ per track (d=3 state dim)	< 2 ms	< 1 ms	Weight optimization ω per track
Centrality	$O(N^3)$ exact, $O(N^2)$ approximate	< 5 ms	< 3 ms (power iteration)	Convergence for near-symmetric graphs